
Hardware-Aware Reformulation of CNNs for Efficient Execution on Specialized AI Hardware: A Case Study on NVIDIA Tensor Cores

Anonymous Authors¹

Abstract

Convolutional Neural Networks (CNNs) are central to modern AI, but their performance is often limited by hardware constraints. NVIDIA Tensor Cores, for instance, require input channels to be multiples of 8 and sometimes 512 for efficient execution (NVIDIA Corporation, 2023). Traditional approaches address such alignment using zero-padding, which can be inefficient. In this work, we present a first-step, hardware-aware reformulation of CNN computations using rewrite rules, restructuring the underlying math to satisfy hardware alignment entirely **post-training** without modifying network weights. While our current implementation focuses on a single transformation for Tensor Cores, this approach is generalizable, laying the foundation to explore additional transformations for other accelerators, including AMD GPUs. This study represents an initial step toward *semantic tuning*, a systematic, hardware-aware optimization strategy for efficient deployment of CNN models on specialized AI hardware.

1. Introduction

Convolutional Neural Networks (CNNs) are a foundational building block of modern deep learning systems, underpinning applications ranging from computer vision to speech and scientific computing. At a high level, CNNs apply learnable filters to structured input tensors in order to extract hierarchical feature representations. For a standard convolution, the input tensor has shape $(N, C_{\text{in}}, H[, W, D])$, the filter weights have shape $(C_{\text{out}}, C_{\text{in}}, K_1[, K_2, \dots, K_n])$, and the resulting output tensor has shape $(N, C_{\text{out}}, H_{\text{out}}[, W_{\text{out}}, D_{\text{out}}])$. This formulation uses the NCHW layout; alternate layout is NHWC.

While this mathematical formulation is hardware-agnostic, the efficient execution of CNNs on modern AI accelerators is strongly influenced by hardware-specific constraints. Specialized units such as NVIDIA Tensor Cores and analogous matrix-multiply engines in other accelerators impose alignment and tiling requirements on tensor dimensions, most

notably that certain dimensions (e.g., channel counts) be multiples of fixed factors such as eight or a higher count like 512 (NVIDIA Corporation, 2023). When these constraints are not satisfied, libraries either fall back to less efficient execution paths or require auxiliary modifications such as zero padding, which introduce redundant computation and memory overhead.

It is possible to address these constraints during network design or training, for example by manually choosing channel dimensions or retraining models with padded tensors. In contrast, this paper explores a complementary and largely underexplored direction: post-training reformulation of CNN computations. Rather than modifying the network architecture or retraining weights, we rewrite the underlying convolutional mathematics to produce an equivalent formulation that satisfies hardware alignment and channel count requirements by construction.

Specifically, we present an initial hardware-aware transformation that reshapes and reinterprets the convolutional computation—through **width folding for input in NHWC format (alternatively height folding for NCHW format)** and **structured filter expansion**—so that the resulting tensors conform to accelerator alignment constraints while preserving exact numerical equivalence to the original CNN. This transformation eliminates the need for zero padding and does not alter the learned parameters or outputs of the network. Although the paper focuses on a single transformation motivated by NVIDIA Tensor Core constraints, the broader goal is to establish a foundation for a family of such rewrite rules that can be systematically derived and applied across different accelerator architectures.

Viewed through this lens, CNN execution can be treated as a compilation problem, where mathematically equivalent reformulations are selected to best match the capabilities and constraints of the target hardware. This work represents a first step toward such a hardware-aware compilation framework for CNNs, demonstrating that non-trivial accelerator constraints can be addressed purely through post-training mathematical rewrites.

To motivate the need for hardware-aware filter alignment, we survey several widely used CNN architectures and their

055
 056 *Table 1.* First-layer channels in popular CNN architectures
 057

Network	Input Channels
AlexNet	3
VGG16	3
ResNet-18	3
ResNet-50	3
GoogLeNet	3
MobileNetV2	3

first-layer channel dimensions. As shown in Table 1, most models processing RGB images have an input channel count of 3, which is not a multiple of 8, limiting the efficiency on NVIDIA Tensor Cores. Furthermore, many models have a first-layer filter count below 512, which is the recommended number mentioned in NVIDIA (NVIDIA Corporation, 2023). In such cases, our method can also increase the channel dimensions, ensuring better alignment with hardware requirements while preserving network functionality.

The paper makes the following contributions:

- Introduce *semantic reformulation* as a hardware-aware optimization paradigm.
- Present *width folding*, a semantics-preserving transformation that increases effective channel dimensionality without padding.
- Provide a formal correctness proof.
- Show how the transformation fits naturally as an MLIR compiler pass.
- Demonstrate generalization to GEMM via 1×1 convolutions.
- Report up to $3\times$ speedups over cuDNN/TensorRT on NVIDIA A100.

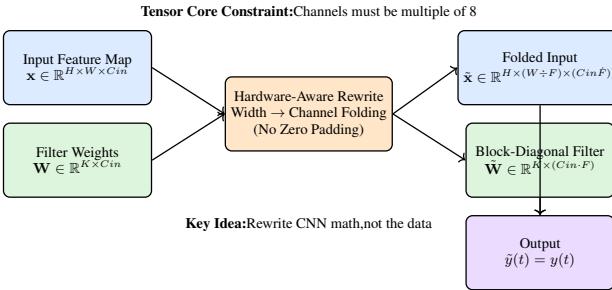


Figure 1. Hardware-aware reformulation of a CNN via width-to-channel folding. The rewrite preserves exact semantics while satisfying Tensor Core alignment constraints without zero padding or retraining.

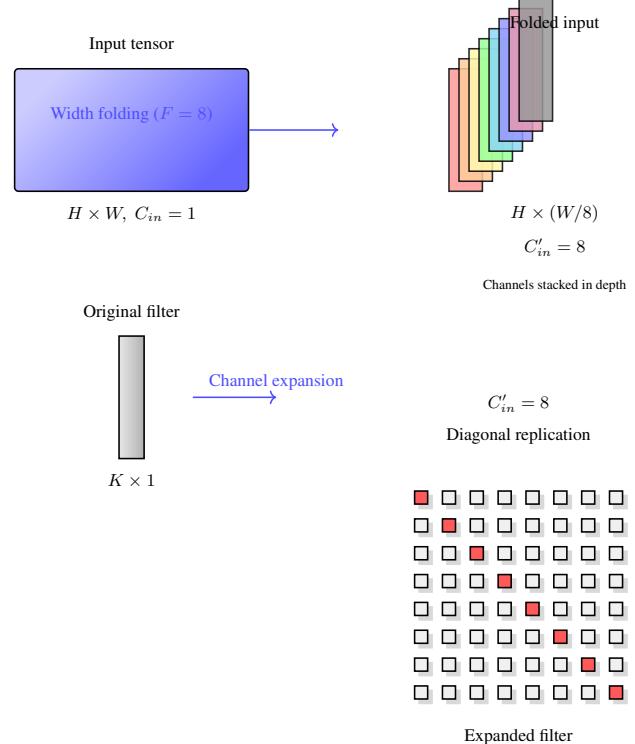


Figure 2. Semantic-preserving CNN reformulation via width folding. The input width is partitioned into $F = 8$ interleaved slices and stacked along the channel dimension, increasing the effective number of input channels without altering spatial height. The original $K \times 1$ convolution kernel is replicated along the main diagonal of the expanded filter matrix, ensuring independent convolution of each folded slice. This post-training transformation preserves exact convolution semantics while aligning channel dimensions with accelerator constraints.

2. Width Folding Transformation

This section describes the width folding transformation, a post-training, semantics-preserving reformulation that increases the effective channel dimension to satisfy hardware alignment constraints while leaving the learned model parameters unchanged.

Given an input tensor

$$\mathbf{x} \in \mathbb{R}^{H \times W \times C_{in}}, \quad C_{in} = 1,$$

and a 1-D convolution kernel

$$\mathbf{W}_f \in \mathbb{R}^{K \times 1},$$

we fold the width dimension by a factor F such that W is divisible by F . The transformation produces an equivalent convolution with

$$C'_{in} = F, \quad W' = W/F,$$

while preserving exact convolution semantics.

```

110
111 Algorithm 1 Width-Folding Transformation for Convolution
112
113 Require: Input tensor  $X \in \mathbb{R}^{B \times H \times W \times C_{\text{in}}}$ ,
114     Filter tensor  $W_f \in \mathbb{R}^{K_H \times K_W \times C_{\text{in}} \times C_{\text{out}}}$ ,
115     Bias  $b \in \mathbb{R}^{C_{\text{out}}}$ ,
116     Folding factor  $F$ 
117 Ensure: Transformed tensors  $(X_f, W'_f, b')$  or fallback
118 1: if  $W \bmod F \neq 0$  or  $C_{\text{in}} \neq 1$  then
119 2:   return  $(X, W_f, b)$  {Fallback: width not divisible or
120     unsupported channels}
121 3: end if
122 4: Define  $X_f \in \mathbb{R}^{B \times H \times (W/F) \times F}$ 
123 5: for  $b\_idx = 0$  to  $B - 1$  do
124 6:   for  $h = 0$  to  $H - 1$  do
125 7:     for  $w' = 0$  to  $(W/F) - 1$  do
126 8:       for  $f = 0$  to  $F - 1$  do
127 9:          $X_f[b\_idx, h, w', f] \leftarrow X[b\_idx, h, F \cdot w' +$ 
128            $f, 0]$ 
129 10:      end for
130 11:    end for
131 12:  end for
132 13: end for
133 14: Define  $W'_f \in \mathbb{R}^{K_H \times K_W \times F \times (F \cdot C_{\text{out}})}$ , initialize to zero
134 15: for  $f = 0$  to  $F - 1$  do
135 16:    $W'_f[:, :, f, f \cdot C_{\text{out}} : (f + 1) \cdot C_{\text{out}}] \leftarrow W_f[:, :, 0, :]$ 
136 17: end for
137 18: Define  $b' \in \mathbb{R}^{F \cdot C_{\text{out}}}$ 
138 19: for  $f = 0$  to  $F - 1$  do
139 20:    $b'[f \cdot C_{\text{out}} : (f + 1) \cdot C_{\text{out}}] \leftarrow b$ 
140 21: end for
141 22: return  $(X_f, W'_f, b')$ 

```

2.1. Input Width Folding

The width folding operation partitions the width dimension into F interleaved slices and stacks them along the channel dimension. Formally, the transformed input tensor

$$X' \in \mathbb{R}^{H \times (W/F) \times F}$$

is defined as

$$X'(h, w', f) = X(h, Fw' + f), \quad f = 0, \dots, F - 1. \quad (1)$$

This operation is a pure re-indexing and does not alter the numerical values of the input tensor.

2.2. Filter Construction

Because convolution is performed only along the height dimension, each width slice is convolved independently. To preserve this behavior after folding, the original filter is replicated across the expanded channel dimension without introducing cross-channel mixing.

Let the original filter be

$$W_f \in \mathbb{R}^{K \times 1}.$$

We construct a new filter tensor

$$W'_f \in \mathbb{R}^{K \times F \times F}$$

as a diagonal replication:

$$W'_f(k, f, f') = \begin{cases} W_f(k), & f = f', \\ 0, & f \neq f'. \end{cases} \quad (2)$$

In implementation terms, this corresponds to allocating a zero-initialized filter tensor and copying the original filter into the diagonal channel blocks. Conceptually, each folded width slice receives an identical copy of the original kernel.

2.3. Bias Construction

The original convolution bias

$$b \in \mathbb{R}$$

is shared across all folded slices. Accordingly, the new bias vector is constructed by replication:

$$b'(f) = b, \quad f = 0, \dots, F - 1. \quad (3)$$

This ensures that each expanded channel applies the same bias as in the original formulation.

2.4. Resulting Convolution

After width folding and filter construction, the transformed convolution operates on

$$X' \in \mathbb{R}^{H \times (W/F) \times F}$$

using the filter W'_f and bias b' . As shown in Section ??, this convolution produces outputs that are exactly equivalent to those of the original network, up to a bijective re-indexing of the width dimension. The algorithm for computing width folded convolution is presented in 1.

3. Width Folding: Mathematical Perspective

Width folding is a structural transformation of convolutional tensors that trades spatial width for additional channels. Given an input tensor $X \in \mathbb{R}^{B \times H \times W \times C_{\text{in}}}$ and a folding factor F , width folding produces $X_f \in \mathbb{R}^{B \times H \times (W/F) \times (C_{\text{in}} \cdot F)}$, via a linear isomorphism

$$X_f[b, h, w', c'] = X[b, h, F \cdot w' + f, c], \quad c' = f \cdot C_{\text{in}} + c,$$

preserving all input information. Conceptually, this is equivalent to reindexing the contraction indices in convolution, where the standard contraction

$$Y[b, h, w, c_{\text{out}}] = \sum_{k_h, k_w, c_{\text{in}}} X[b, h+k_h, w+k_w, c_{\text{in}}] W[k_h, k_w, c_{\text{in}}, c_{\text{out}}]$$

165 becomes, after width folding,

$$166 \quad Y_f[b, h, w', c'_{\text{out}}] = \sum_{k_h, k_w, c'_{\text{in}}} X_f[b, h+k_h, w', c'_{\text{in}}] W'_f[k_h, k_w, c'_{\text{in}}, f]$$

$$167 \quad \text{formed input tensor}$$

$$168$$

$$169$$

170 In linear algebra terms, folding flattens width into channels,
 171 allowing the convolution to be represented as a block-diagonal matrix multiplication, where each block corresponds to one slice along the folded width. This structure is
 172 naturally interpreted as a Kronecker product: the expanded kernel W'_f can be viewed as the Kronecker product of the
 173 original kernel with an identity along the folded width dimension, which aligns well with hardware vectorization and
 174 parallelism.

175 From a category-theoretic viewpoint, tensors are objects
 176 in a monoidal category (e.g., $\text{Vect}_{\mathbb{R}}$) and convolution is
 177 a morphism $X \otimes W \rightarrow Y$. Width folding is a natural
 178 isomorphism of the form

$$179 \quad B \otimes H \otimes W \otimes C_{\text{in}} \cong B \otimes H \otimes W' \otimes (C_{\text{in}} \otimes F),$$

$$180$$

$$181$$

$$182$$

$$183$$

184 reassociating the tensor factors without changing the underlying morphism. Thus, folding is a structure-preserving
 185 transformation: the semantics of convolution remain identical, while the indexing structure changes to facilitate computation.

186 Overall, width folding unifies perspectives from tensor re-
 187 shaping, contraction, block-diagonal / Kronecker structures,
 188 and categorical isomorphisms, providing a concise mathematical framework for this hardware-oriented optimization.

4. Correctness of Width Folding Transformation

191 We consider a 1-D convolution performed exclusively along
 192 the height dimension H . The width dimension W does not
 193 participate in the convolution and is treated as an independent
 194 indexing dimension.

195 Let the input tensor be

$$196 \quad X \in \mathbb{R}^{H \times W \times C_{\text{in}}}, \quad C_{\text{in}} = 1,$$

$$197$$

$$198$$

$$199$$

$$200$$

$$201$$

$$202$$

$$203$$

$$204$$

$$205$$

$$206$$

$$207$$

$$208$$

$$209$$

210 and let the convolution filter be

$$211 \quad W_f \in \mathbb{R}^{K \times 1},$$

$$212$$

$$213$$

$$214$$

$$215$$

$$216$$

$$217$$

$$218$$

$$219$$

220 with bias $b \in \mathbb{R}$. Batch and output channel indices are
 221 omitted for clarity.

222 The original convolution output is given by

$$223 \quad Y(h, w) = \sum_{k=0}^{K-1} W_f(k) X(h+k, w) + b, \quad (4)$$

$$224$$

$$225$$

$$226$$

$$227$$

$$228$$

$$229$$

230 where convolution is performed only along the height dimension.

Width Folding Transformation. Let F be the folding factor and assume W is divisible by F . We define a transformed input tensor

$$X' \in \mathbb{R}^{H \times (W/F) \times F}$$

by re-indexing the width dimension as

$$X'(h, w', f) = X(h, Fw' + f), \quad f = 0, \dots, F-1. \quad (5)$$

This operation folds the width dimension into the channel dimension without modifying spatial data along H .

Filter Expansion. Since the convolution is independent across width indices, the filter must be replicated for each folded slice. We define an expanded filter

$$W'_f \in \mathbb{R}^{K \times F \times F}$$

as a diagonal replication:

$$W'_f(k, f, f') = \begin{cases} W_f(k), & f = f', \\ 0, & f \neq f'. \end{cases} \quad (6)$$

The bias is similarly replicated as $b'(f) = b$.

Transformed Convolution. The output of the transformed convolution is

$$Y'(h, w', f) = \sum_{k=0}^{K-1} \sum_{f'=0}^{F-1} W'_f(k, f, f') X'(h+k, w', f') + b'(f). \quad (7)$$

Substituting Eq. (6) into Eq. (7) removes the channel summation:

$$Y'(h, w', f) = \sum_{k=0}^{K-1} W_f(k) X'(h+k, w', f) + b. \quad (8)$$

Using the folded input definition from Eq. (5), we obtain

$$Y'(h, w', f) = \sum_{k=0}^{K-1} W_f(k) X(h+k, Fw' + f) + b.$$

Let $w = Fw' + f$. Then

$$Y'(h, w', f) = Y(h, w),$$

where $Y(h, w)$ is the output of the original convolution defined in Eq. (4).

Conclusion. The width folding transformation constitutes a bijective re-indexing of the width dimension combined with diagonal filter replication. Since the convolution is performed solely along the height dimension, the transformation preserves the exact numerical output of the original network. Therefore, width folding with channel expansion is *semantics preserving*. \square

220 4.1. N-D Convolutions

221 Our method generalizes to N-D convolutions by folding
 222 dimensions that are not involved in the convolution operation
 223 into the channel dimension. This reparameterization
 224 preserves exact convolution semantics via block-diagonal
 225 kernels and does not rely on kernel separability or approxi-
 226 mation.
 227

228 5. Width-Folding Transformation as an IR 229 Transformation

230 5.1. Motivation

231 Formulating width-folding as a compiler transformation is
 232 feasible and essential because it decouples the mathematical
 233 optimization from any particular library or runtime. By
 234 representing the input reorganization and kernel replication at
 235 the IR level, the compiler can reason about the data layout,
 236 tiling, and vectorization systematically. This enables auto-
 237 matic generation of hardware-efficient code that maximizes
 238 utilization of compute units such as Tensor Cores, while
 239 preserving the semantics of the original convolution. Treat-
 240 ing it as a compiler pass ensures portability, composability
 241 with other optimizations, and the ability to target multiple
 242 backends without manual intervention.
 243

244 5.2. Mathematical Formulation

245 Let the input tensor be $\mathcal{X} \in \mathbb{R}^{B \times H \times W \times C_{\text{in}}}$, and the ker-
 246 nel be $\mathcal{K} \in \mathbb{R}^{K_H \times K_W \times C_{\text{in}} \times C_{\text{out}}}$, producing output $\mathcal{Y} \in$
 247 $\mathbb{R}^{B \times H' \times W' \times C_{\text{out}}}$.

248 The *width-folding* transformation reshapes the input tensor
 249 along the width dimension. F (folding factor) is chosen to
 250 align with Tensor core tile sizes.
 251

$$252 \mathcal{X}_f \in \mathbb{R}^{B \times H \times \frac{W}{F} \times (C_{\text{in}} \cdot F)}, \quad F \in \mathbb{Z}^+ \quad (9)$$

253 The kernel is correspondingly transformed into a diagonal-
 254 blocked form:
 255

$$256 \mathcal{K}_f \in \mathbb{R}^{K_H \times \frac{K_W}{F} \times (C_{\text{in}} \cdot F) \times (C_{\text{out}} \cdot F)}. \quad (10)$$

257 The transformed convolution preserves the original seman-
 258 tics:
 259

$$260 \mathcal{Y} = \text{Conv}(\mathcal{X}, \mathcal{K}) = \text{Reconstruct}(\text{Conv}(\mathcal{X}_f, \mathcal{K}_f)), \quad (11)$$

261 where the reconstruction step involves reshaping the folded
 262 output back to the original width dimension.
 263

264 5.3. Compiler-Level Realization in MLIR

265 Width-folding can be potentially implemented
 266 as a *semantics-preserving IR transformation in*
 267 *MLIR*. The transformation shall be performed over
 268 `linalg.conv.2d_nhwc` or `linalg.matmul` operations, and potentially consist of the following steps:

- 269 1. **Tensor Reshape:** Reshape the input and output tensors
 270 to introduce a blocked channel dimension corresponding
 271 to the folding factor F .
- 272 2. **Affine Reindexing:** Map the original convolution indices
 273 to the folded tensor indices, ensuring that all data
 274 dependencies are preserved.
- 275 3. **Kernel Replication:** Replicate the kernel into a
 276 diagonal-blocked layout that is consistent with the
 277 folded input tensor.

278 The legality of this transformation is straightforward: the
 279 width dimension must be divisible by F , and any padding
 280 must be applied consistently to preserve output shape. A
 281 lightweight cost model can estimate profitability by consider-
 282 ing channel size, tensor core tile alignment, and arithmetic
 283 intensity. The transformation is fully composable with other
 284 MLIR passes such as tiling, vectorization, and lowering to
 285 CUDA or ROCm backends. Width folding thus provides
 286 a mathematically sound, hardware-aware optimization that
 287 can significantly improve performance of deep learning ker-
 288 nels in production compilers.

289 6. Generalization to GEMM via 1×1 290 Convolution

291 While width folding is introduced in the context of convo-
 292 lutional operators, the underlying idea extends naturally to
 293 general matrix multiplication (GEMM). This follows from
 294 the well-known equivalence between GEMM and 1×1
 295 convolution under appropriate reshaping. We describe this
 296 equivalence and show how width folding applies directly to
 297 GEMM through this formulation.

298 6.1. GEMM as a 1×1 Convolution

299 Consider a matrix multiplication

$$300 C = AB,$$

301 where $A \in \mathbb{R}^{M \times K}$ and $B \in \mathbb{R}^{K \times N}$.

302 We reshape matrix A into a tensor

$$303 X \in \mathbb{R}^{H \times W \times C_{\text{in}}},$$

304 by choosing

$$305 H = M, \quad W = 1, \quad C_{\text{in}} = K,$$

275 and interpreting each row of A as a spatial position with K
 276 channels.

277 Similarly, matrix B is reshaped into a 1×1 convolution
 278 kernel

$$279 \quad W_f \in \mathbb{R}^{1 \times 1 \times C_{\text{in}} \times C_{\text{out}}},$$

280 with $C_{\text{out}} = N$.

281 Applying a 1×1 convolution produces an output tensor

$$282 \quad Y \in \mathbb{R}^{H \times 1 \times C_{\text{out}}},$$

286 which is equivalent to the GEMM result C after reshaping.

287 This construction is algebraically exact and introduces no
 288 approximation.

290 291 6.2. Applying Width Folding to GEMM

292 In this formulation, the channel dimension corresponds to
 293 the reduction dimension K of the matrix multiplication.
 294 When K is small or poorly aligned with hardware tile sizes,
 295 the resulting computation underutilizes specialized matrix-
 296 multiply units.

297 Width folding can be applied by reindexing an auxiliary spatial
 298 dimension and redistributing it into the channel dimension.
 299 Concretely, we introduce a synthetic width dimension
 300 W and fold it into channels:

$$301 \quad X \in \mathbb{R}^{H \times W \times 1} \rightarrow X' \in \mathbb{R}^{H \times (W/F) \times F}.$$

302 The corresponding 1×1 kernel is replicated into a block-
 303 diagonal form, exactly as in the convolutional case. The
 304 resulting operator remains a 1×1 convolution and is there-
 305 fore equivalent to a GEMM with expanded effective channel
 306 dimensionality.

307 This equivalence shows that width folding is not specific
 308 to convolutional layers, but applies more broadly to linear
 309 operators expressed as tensor contractions. From a compiler
 310 perspective, GEMM and convolution differ only in indexing
 311 structure, and both can benefit from operator-level reformu-
 312 lation. This observation further supports expressing width
 313 folding as a general compiler transformation rather than a
 314 domain-specific optimization.

315 316 7. Potential Optimizations for Sparsity and 317 Quanitzation

318 Efficient exploitation of sparsity is crucial for both memory
 319 savings and computational acceleration. The block-diagonal
 320 filter structure introduces structured sparsity that can be
 321 leveraged in multiple ways. Only the diagonal blocks of the
 322 filters contain non-zero values, while off-diagonal blocks
 323 are zero. This structured sparsity allows storing just the
 324 non-zero blocks in memory as sparse tensors, significantly

325 reducing memory footprint. For large networks, this can
 326 lead to substantial savings, especially in embedded or GPU-
 327 constrained environments. Several modern deep learning
 328 frameworks and custom CUDA kernels can exploit this by
 329 performing sparse matrix multiplications, reducing the num-
 330 ber of operations and increasing throughput. Popular deep
 331 learning frameworks such as PyTorch and TensorFlow pro-
 332 vide built-in support for grouped convolutions. By mapping
 333 each diagonal block to a group, the block-diagonal convolu-
 334 tion can be implemented efficiently without modifying the
 335 core framework.

336 Structured sparsity pairs naturally with mixed-precision
 337 quantization (FP16/INT8). By representing both weights
 338 and activations in lower precision, memory bandwidth is
 339 reduced, and Tensor Cores can achieve higher throughput.
 340 Combining quantization with block-diagonal sparsity en-
 341 sures that only essential computations are performed at
 342 high speed, improving overall efficiency while maintain-
 343 ing model accuracy.

344 8. Results

345 We implemented the proposed widthfolding transformation
 346 in both TensorFlow and TensorRT, validating it against func-
 347 tionally equivalent convolution operations. All experiments
 348 and ablation studies were conducted on NVIDIA A100
 349 GPUs using proprietary CNN models. Due to intellectual
 350 property and data-sharing constraints, detailed experimental
 351 configurations, model architectures, and raw performance
 352 measurements cannot be publicly disclosed. However, a
 353 reference implementation in Python using TensorFlow API
 354 is provided in Appendix A, allowing reproducibility of the
 355 transformation on synthetic or publicly available models.

356 To provide context for performance, the complete set of
 357 empirical results is summarized in the accompanying patent
 358 application (Bikshandi & Seberino, 2024). These results
 359 demonstrate a minimum of 3x improvement over baseline
 360 implementations (vanilla TensorRT C++ API (NVIDIA,
 361 2023)) on A100-class hardware. The observed gains are
 362 primarily attributable to alignment-aware reformulation of
 363 convolution operations, which improves Tensor Core utiliza-
 364 tion compared to conventional zero-padding strategies.

365 FP16 precision was used in all experiments. It should
 366 be noted that TensorRT does not natively support Tensor
 367 Core operations with FP16 input when the number of input
 368 channels is not a multiple of eight. It was using a sub-
 369 optimal convolution kernel for this case. While H100 and
 370 B-series GPUs were unavailable during experimentation,
 371 the method is expected to generalize to newer architectures
 372 (H200/H800/B200), and re-evaluation on these platforms
 373 is planned contingent on hardware access. Importantly,
 374 TensorRT is already highly optimized; achieving 3x per-

330 performance improvement indicates a fundamentally new al-
 331 gorithmic approach, improved hardware utilization, and
 332 domain-specific optimization beyond generic kernels.

333 Although cuDNN (Chetlur et al., 2014), another popu-
 334 lar deep learning framework, exposes FP16 convolution at
 335 the API level, convolutions with very small input channel
 336 counts—most notably $C_{\text{in}} = 1$ —are not allowed by default
 337 due to reason that tensor Cores require matrix dimensions
 338 that are multiples of 8. For such cases, the responsibility
 339 is pushed to the user rather than the library and users must
 340 explicitly pad the channel dimension with zeros or use FP32
 341 variant. Zero padding increases memory traffic and per-
 342 forms wasted computation on artificial data, while falling
 343 back to FP32 introduces FP16–FP32–FP16 conversion over-
 344 head and executes on CUDA cores (not Tensor cores). In
 345 both cases, Tensor Cores remain underutilized, despite their
 346 theoretical ability to provide up to an 8× throughput im-
 347 provement for FP16 operations.
 348

349 The above facts reflect the design philosophy of TensorRT
 350 and cuDNN as a general-purpose kernel library rather than
 351 a system that invents new mathematical formulations. They
 352 faithfully implement the operation given; they do not reinter-
 353 pret tensor dimensions or apply semantics-changing trans-
 354 formations to satisfy hardware alignment requirements. In
 355 contrast, our approach introduces a semantics-preserving
 356 mathematical transformation that increases the effective
 357 channel dimension without zero padding or precision
 358 up/down-conversion. This reparameterization makes the
 359 convolution natively compatible with FP16 Tensor Core
 360 kernels, enabling high utilization and substantially higher
 361 performance than existing cuDNN or TensorRT implemen-
 362 tations for low-channel-count convolutions.
 363

364 The proposed method applies to any convolutional layer, not
 365 just the first layer, and can also be extended to linear layers
 366 (i.e., GEMM) as noted in Section 6. This extension implies
 367 that the transformation could accelerate GEMM operations
 368 in cuBLAS (especially, those involving *tall skinny matrices*).
 369 A 3× speedup over cuBLAS is a potential break through in
 370 compiler optimization.

371 The observed results are consistent with publicly available
 372 NVIDIA documentation (NVIDIA Corporation, 2023), pro-
 373 viding further confidence in the correctness and potential
 374 impact of the method. Finally, this transformation can be
 375 applied even during tranining process to realize similar
 376 speedups and tensorcore utilization.
 377

378 9. Related Work

380 Optimizing CNNs for hardware accelerators has been an
 381 active area of research. Our approach draws inspiration
 382 from several exisiting techniques, combining channel expan-
 383 sion and block-diagonal sparsity to fully exploit hardware
 384

throughput.

9.1. Relation to Other CNN Variants

9.1.1. GROUPED CONVOLUTIONS

In grouped convolutions (Krizhevsky et al., 2012), input channels are divided into separate groups, and each group is convolved independently with a corresponding set of filters. Our block-diagonal transformation naturally induces a similar structure: each diagonal block in the filter corresponds to a “group” that processes a subset of the input channels. However, unlike standard grouped convolutions, our approach preserves the original filter values within each block and replicates them in a controlled manner to align with Tensor Core hardware. This ensures maximum hardware utilization without changing the semantic meaning of the original convolution.

9.1.2. DEPTHWISE AND POINTWISE CONVOLUTIONS

Depthwise convolutions (Howard et al., 2017) apply a single filter per input channel, significantly reducing computation but also limiting feature interactions across channels. Pointwise (1×1) convolutions are then used to combine features across channels. The block-diagonal filter transformation resembles a hybrid of these approaches: each block is multi-channel (not single-channel as in depthwise) and operates on a subset of input channels, while off-diagonal blocks are zeroed to create structured sparsity. This design maintains full multi-channel processing within blocks while still enabling computational efficiency similar to depthwise convolutions. Unlike traditional depthwise or pointwise convolutions, our method explicitly targets hardware alignment for Tensor Cores.

9.1.3. CHANNEL EXPANSION / 1×1 CONVOLUTIONS

Channel expansion via 1×1 convolutions is commonly used to increase the number of output channels and enable more expressive representations. In our approach, the block-diagonal transformation can be interpreted as a structured channel expansion: the original output channels are replicated across diagonal blocks to create a padded, hardware-aligned output tensor. This provides the benefits of channel expansion—higher representational capacity—without introducing additional learnable parameters beyond the original filter. The key distinction is that our expansion is carefully organized to match hardware constraints, which is not considered in conventional 1×1 convolutions.

Similar strategies have been explored in low-rank filter approximations (Jaderberg et al., 2014) and kernel tiling optimizations (Chetlur et al., 2014), but our method maintains the original filter semantics while providing full Tensor Core alignment.

385
386
387
388
389
390
391
392
393
394
395
396
397

9.2. Difference from Channel Zero-Padding

A simple approach to align channels for hardware is *zero-padding*: if a layer has 1 input channel and the hardware prefers multiples of 8, 7 zero channels are added. While this aligns the total channels, these extra channels carry no information, wasting compute on irrelevant data. In contrast, *width-folding* avoids adding input channels, expanding only the (small) weight tensor. It creates *block-diagonal filters*, introducing structured sparsity that frameworks can exploit for efficient computation. Moreover, this pattern naturally generalizes to higher-dimensional and grouped convolutions.

9.3. Contrast with Traditional Compiler Transformations

Width folding differs fundamentally from traditional compiler transformations applied in deep learning systems. Traditional compiler optimizations—such as operator fusion, loop tiling, layout reordering, vectorization, and kernel selection—preserve the *high-level mathematical definition* of an operator and optimize *how* that computation is executed on hardware. These transformations act on scheduling, memory access patterns, or code generation, but do not alter the underlying convolution formulation.

In contrast, width folding operates at the level of the *mathematical operator itself*. The transformation explicitly rewrites the convolution into an equivalent form by re-indexing tensor dimensions and restructuring filter parameters. This reformulation changes the apparent tensor shapes and filter structure seen by the compiler, while preserving exact input–output semantics by construction. As a result, hardware alignment constraints (e.g., channel dimensions being multiples of a fixed factor) are satisfied without relying on padding or special-case kernel implementations.

Finally, while compiler transformations are generally opaque and hardware-specific, width folding is expressed as an explicit, semantics-preserving rewrite rule. This makes the transformation analyzable, provably correct, and amenable to future automation.

Width folding is a *post-training, pre-compilation* transformation: it modifies the trained model itself before it is handed to the compiler or runtime. This enables downstream compilers to treat the transformed model as a standard convolution that naturally maps to optimized hardware kernels. Rather than replacing compiler optimizations, width folding complements them by expanding the space of hardware-friendly formulations that compilers can further optimize.

435
436
437
438
439

10. Conclusion

We presented a transformation for CNN filters that aligns input and output channels with NVIDIA Tensor Core requirements. The method maintains the original filter structure, generalizes to n-D convolutions, and exploits sparsity for memory and computational efficiency. This approach bridges hardware-aware optimizations with traditional CNN design principles, providing a practical and scalable solution for high-performance deep learning deployments.

10.1. Future Work

Several promising research directions follow naturally. First, we plan to explore a broader class of *structure-preserving transformations* beyond block-diagonal reformulations. These include alternative spatial–channel folding strategies, hierarchical blocking, mixed-radix reshaping schemes, sparsity aware folding and dilation rewriting. Such transformations could enable efficient utilization of hardware units with strict alignment, vectorization, or tiling requirements. Also, it is interesting to study if the inverse (i.e. channel-to-space) is possible mathematically.

Second, we plan to investigate transformations tailored to other GPU architectures, including AMD GPUs, where waveform sizes, memory coalescing rules, and matrix instruction formats differ substantially. Developing a unified transformation framework that adapts this idea to multiple specialized hardware as well as general-purpose CPU (e.g. alignment requirement for SIMD or cache usage ([oneAPI Deep Neural Network Library \(oneDNN\) Team, 2025](#))) remains an open and important problem.

Third, we aim to extend these techniques to other operators commonly used in modern architectures, such as grouped convolutions, depthwise separable convolutions, attention mechanisms ([Vaswani et al., 2017](#)), and recurrent layers ([Hochreiter & Schmidhuber, 1997](#)).

Finally, we envision integrating these transformations into MLIR ([Lattner et al., 2021](#)) framework. More broadly, we introduce *semantic tuning* as a post-training optimization paradigm in which the mathematical formulation of a neural network is rewritten to better match hardware execution constraints while preserving the exact input–output semantics of the original model. This shifts the compiler optimizations from hardware-specific kernel optimization toward *semantically equivalent operator transformations*, achieving performance portability.

References

- Bikshandi, G. and Seberino, C. Efficient execution of machine learning models on specialized hardware, 2024.
 URL <https://patents.google.com/patent/>

- 440 WO2024197121A1/en. International publication date:
441 2024-09-26; prior art keywords include machine learning
442 model configuration and specialized computing devices.
443
- 444 Chetlur, S. et al. cudnn: Efficient primitives for deep
445 learning. <https://arxiv.org/abs/1410.0759>,
446 2014. arXiv:1410.0759.
- 447 Hochreiter, S. and Schmidhuber, J. Long short-term memory.
448 *Neural Computation*, 9(8):1735–1780, 1997.
449
- 450 Howard, A. G. et al. Mobilenets: Efficient convolu-
451 tional neural networks for mobile vision applica-
452 tions. <https://arxiv.org/abs/1704.04861>,
453 2017. arXiv:1704.04861.
- 454 Jaderberg, M., Vedaldi, A., and Zisserman, A. Speeding up
455 convolutional neural networks with low rank expansions.
456 In *Proceedings of the British Machine Vision Conference*
457 (*BMVC*), 2014.
- 458 Krizhevsky, A., Sutskever, I., and Hinton, G. E. Imagenet
459 classification with deep convolutional neural networks.
460 In *Advances in Neural Information Processing Systems*
461 (*NeurIPS*), 2012.
- 462 Lattner, C., Pienaar, J., Amini, M., Bondhugula, U., Cohen,
463 A., Davis, A., Pienaar, J. A., Riddle, R., Shpeisman,
464 T., Vasilache, N., and Zinenko, O. Mlir: A compiler
465 infrastructure for the end of moore’s law. In *Proceed-
466 ings of the IEEE/ACM International Symposium on Code
467 Generation and Optimization (CGO)*, 2021.
- 468 NVIDIA. TensorRT: High-performance deep learning infer-
469 ence optimization. <https://developer.nvidia.com/tensorrt>, 2023.
- 470 NVIDIA Corporation. Optimizing convolutional
471 layers user guide. <https://docs.nvidia.com/deeplearning/performance/pdf/Optimizing-Convolutional-Layers-User-Guide.pdf>, 2023. Accessed: 2025-01-XX.
- 472 oneAPI Deep Neural Network Library (oneDNN) Team.
473 *oneDNN Developer Guide and Reference, Version
474 2025.1*. oneAPI / Intel Corporation, 2025. URL <https://www.intel.com/content/www/us/en/docs/oneapi/programming-guide/2024-1/intel-oneapi-deep-neural-network-library-onednn.html>. Accessed: 2026-01-04.
- 475 Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones,
476 L., Gomez, A. N., Kaiser, L., and Polosukhin, I. Atten-
477 tion is all you need. In *Advances in Neural Information
478 Processing Systems (NeurIPS)*, 2017.
- 479
- 480
- 481
- 482
- 483
- 484
- 485
- 486
- 487
- 488
- 489
- 490
- 491
- 492
- 493
- 494

A. Example Code using Tensorflow CNN (Channels-Last)

496 This Tensorflow implementation demonstrates a 2-D block-diagonal filter transformation where the input is in **channels-last
497 format** ('[B, H, W, C;n']'), the *W* dimension is folded, and the convolution is compatible with NVIDIA Tensor Cores.
498

listings xcolor

Listing 1. Width-folding CNN transformation in TensorFlow (channels-last).

```
502
5031 import tensorflow as tf
5042 import numpy as np
5053
5064 # -----
5075 # Parameters
5086 # -----
5097 B, H, W = 1, 32, 64      # batch, height, width
5108 K = 5                  # kernel size along H
5119 F = 8                  # width folding factor
51210 Cout = 1              # output channels
51311
51412 assert W % F == 0
51513
51614 # -----
51715 # Input tensor (NHWC)
51816 #
51917 x = tf.random.normal((B, H, W, 1))
52018
52119 # -----
52220 # Original filter + bias
52321 # Conv along H only -> kernel (K,1)
52422 #
52523
52624 # -----
52725 # Original convolution
52826 #
52927
53028 y_orig = tf.nn.conv2d(
53129     x,
53230     filterVal,
53331     strides=[1, 1, 1, 1],
53432     padding="VALID",
53533     data_format="NHWC"
53634)
53735 y_orig = tf.nn.bias_add(y_orig, biasVal)
53836
53937 # -----
54038 # Width folding: W -> W/F, Cin -> F
54139 #
54240 # (B, H, W, 1) -> (B, H, W/F, F)
54341 x_folded = tf.reshape(x, (B, H, W // F, F))
54442
54543 # -----
54644 # Build diagonal filter
54745 # (K,1,1,Cout) -> (K,1,F,F*Cout)
54846 #
54947
55048 filterValNew = np.zeros((K, 1, F, F * Cout), dtype=np.float32)
55149
55250 for f in range(F):
55351     filterValNew[:, :, f, f*Cout:(f+1)*Cout] = np.squeeze(filterVal.numpy(), axis=-1)
55452
55553 filterValNew = tf.constant(filterValNew)
55654
55755 # -----
55856 # Bias replication
55957
```

```

550 57# -----
551 58biasValNew = tf.tile(biasVal, [F])
552 59
553 60# -----
554 61# Folded convolution
555 62# -----
556 63y_folded = tf.nn.conv2d(
557 64    x_folded,
558 65    filterValNew,
559 66    strides=[1, 1, 1, 1],    # stride only along H
560 67    padding="VALID",
561 68    data_format="NHWC"
562 69)
563 70y_folded = tf.nn.bias_add(y_folded, biasValNew)
564 71
565 72# -----
566 73# Reconstruct original layout
567 74# (B, H', W/F, F) -> (B, H', W)
568 75# -----
569 76y_reconstructed = tf.reshape(
570 77    y_folded,
571 78    (B, y_folded.shape[1], W)
572 79)
573 80
574 81# -----
575 82# Verification
576 83# -----
577 84max_error = tf.reduce_max(tf.abs(np.squeeze(y_orig, axis=-1) - y_reconstructed))
578 85print("Max absolute error:", max_error.numpy())
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604

```